

Hardware-Constrained Architecture Search for Lane-Change Intention Prediction from Vehicle Time Series

Sepehr Mohammady, [TODO: co-authors, order], [TODO: supervisor]

Abstract—Predicting a driver’s lane-change intention early enough to act on it is a time-series classification problem with hard deployment limits: the model has to run on automotive-grade microcontrollers within tight memory and latency budgets. We apply constrained neural architecture search to the Lane Change Intention Recognition driving dataset [3] (windows of 50 samples over 31 signals) and obtain compact one-dimensional convolutional networks for three tasks: three-class intention classification and time-to-lane-change regression for left and right maneuvers. [TODO: results: accuracy, RMSE, model size, latency on STM32] The searched models improve on the published baseline (Transformer, 0.51 s time-to-lane-change RMSE, deployed in FP32 [1]) while using [TODO: X] $\times$ less memory, and we report measured, not estimated, footprint and latency on the same STM32 targets after int8 quantization.

Index Terms—Neural architecture search, lane-change intention, time-series classification, TinyML, embedded inference, STM32.

I. INTRODUCTION

[TODO: Motivation: intention prediction as advance warning for ADAS; why on-device inference (latency, privacy, cost, no connectivity assumptions); the gap: published DMIR models target accuracy, not deployment budgets.]

Contributions:

- 1) a constrained search space and objective tuned to one-dimensional driving signals under microcontroller budgets;
- 2) models that beat the published DMIR baseline on all three tasks at a fraction of its footprint [TODO: verify once results exist];
- 3) measured deployment numbers (flash, RAM, latency) on [TODO: STM32 part], with the full pipeline released as open source.

II. RELATED WORK

A. Lane-change intention prediction

[TODO: The LC dataset and framework baseline [1] (TTL regression, Transformer RMSE 0.51 s, FP32 deployment on STM32H7B3/F401); earlier HighD-based TTL work (RMSE 0.63–0.7); MicroNAS for time series and TinyTNAS as closest NAS relatives.]

The authors are with the ELIOS Lab, Department of Electrical, Electronic and Telecommunications Engineering and Naval Architecture (DITEN), University of Genoa, Genoa, Italy.

This work was supported by the SYNERGIES project [TODO: grant number and official acknowledgment text].

B. Constrained architecture search for microcontrollers

[TODO: muNAS [2], MCUNet/TinyNAS, MicroNets; what transfers to 1D time series and what does not.]

III. DATA

We use the Lane Change Intention Recognition dataset [3] (50 human drivers in the CARLA simulator, 10 Hz, over 3,400 annotated lane changes; MIT license), prepared as fixed windows of $T=50$ samples (5 s) over $F=31$ channels [TODO: confirm channel order and whether channel 31 is fileTime — drop it if so; official feature count is 30]. Classification uses balanced splits of 94,620 / 16,332 / 19,722 windows (train/val/test) over three classes: no intention, right lane change, left lane change. The regression tasks map a window to the time until the lane change, $t \in [0, 4]$ s discretized at 0.1 s, with 39,313 / 7,742 / 8,097 (LCL) and 33,201 / 5,828 / 7,073 (LCR) windows.

The splits are user-independent: we verified by matching raw per-driver session logs against the prepared windows (correlation over 50-sample windows) that the validation and test drivers follow the protocol of [1] with no window leakage across splits. The test splits contain isolated spikes (up to 5×10^6) on the right/left curvatureDx channel pair — a numerical-derivative artefact absent from training; we clip validation and test features to the per-channel training range and leave training data untouched.

IV. METHOD

A. Search space

The search space is built from one-dimensional convolutional blocks: standard, depthwise and 1×1 convolutions with kernel sizes in $\{3, 5, 7\}$ and channel counts up to 128, optional stride-2 and pooling, blocks connected serially or with residual branches, and a head of one to three fully-connected layers. The operator set is restricted to layers the ST Edge AI compiler maps efficiently. [TODO: state the exact ranges once frozen.]

B. Constraints and search algorithm

We use aging evolution: a population of architectures is sampled, the fittest of a random subset is mutated by a network morphism, its child is trained, and the oldest member is retired. Fitness combines four objectives—validation error, peak RAM, model size and multiply–accumulate count—by the random scalarisation of [2]: each is normalised by a user budget

and divided by a random weight resampled every round, so a candidate exceeding any budget is penalised first. Budgets were set from the STM32H7B3I-DK, the deployment platform of the published baseline, so the on-device comparison is like for like. Each task ran 150 rounds with a population of 50 and a sample size of 15; all four searches completed overnight on a single laptop GPU (RTX 5070), against thousands of GPU-days for early reinforcement-learning and evolutionary NAS [5], [6]—a different-task comparison offered as context only, but one that follows from searching directly in the microcontroller-scale regime, where every candidate trains in minutes.

C. Discovered architecture

Across tasks the search independently arrives at the same microcontroller-friendly recipe: 1×1 convolutions to re-shape the channel dimension, depthwise convolutions for temporal structure, aggressive early downsampling, and a small flattened head (98–308 features against the reference’s 6400). Topology, however, follows the task: the classifiers are essentially sequential, whereas the regression model uses two parallel branches merged by an addition—the search space permits branches everywhere, but only regression exploited them. The classification model reduces the sequence $50 \rightarrow 25 \rightarrow 13 \rightarrow 7 \rightarrow 4$ before its head:

$50 \times 31 \rightarrow \text{MaxPool} \rightarrow \text{DWConv}(s2) \rightarrow \text{Conv}_{1 \times 1}(77) \rightarrow \text{DWConv} \rightarrow \text{DWConv}(s2) \rightarrow \text{Conv}_{1 \times 1} \rightarrow \text{Add} \rightarrow \text{AvgPool} \rightarrow \text{FC}(223) \rightarrow \text{FC}(23) \rightarrow \text{FC}(3)$

This explains the footprint gap quantitatively. Both our model and the reference CNN spend most of their flash in the fully-connected head, so the gain does not come from avoiding a large head; it comes from what enters it. The reference holds full temporal resolution and flattens a 50×128 map into 6400 features, giving a 64×6400 layer of 1.64 MB ($\sim 95\%$ of the model). Ours pools first and flattens only $4 \times 77 = 308$ features, so it affords a *wider* head (223 units) in a ~ 280 KB layer—a $6 \times$ reduction. The same downsampling cuts the convolutional work, yielding the $12.4 \times$ fewer MACs and the $9.2 \times$ measured speed-up of Table II.

V. EXPERIMENTS

A. Setup

[TODO: training schedule, hardware, seeds/repeats, hyper-parameters; all runs logged to the released experiments.jsonl.]

B. Results

Table I reports test-set metrics. All searched-model numbers are our own test-set evaluation of the saved models, from searches run to their full 150-round budget. [TODO: add peak-RAM, MACs, and measured on-device latency columns from ST Edge AI; consider re-running with an all-models save policy for a guaranteed-optimal front (see notes).]

For classification the searched model is both smaller and more accurate than the 441 k reference CNN (92.1% at 84 k parameters; a 8 k model still reaches 91.3%). Regression beats the published single-TTLC SOTA (RMSE 0.51) at $2\text{--}3 \times$ fewer

TABLE I
TEST-SET COMPARISON. THE INTERNAL REFERENCE MODELS (UNPUBLISHED) WERE EVALUATED ON THE SAME TEST SET; CLASSIFICATION ACCURACY IS OUR OWN EVALUATION, THE TRANSFORMER RMSE VALUES ARE AS REPORTED. SIZE IS THE PARAMETER COUNT.

Model	Params	Acc. (%)	LCR	LCL
Reference CNN (cls)	441 k	91.5	—	—
Reference Transf. (reg, RMSE)	333/49 k	—	0.42	0.44
Our DSCNN baseline	10 k	91.5	0.32 [†]	0.33 [†]
Ours searched (cls / reg MAE)	84 k	92.1	0.29 [†]	0.33 [†]
Ours compact	8 k	91.3	—	—
Ours, no turn-signal	21 k	91.1	—	—

[†]MAE (our search optimizes MAE); reference regression is reported as RMSE. Our LCR RMSE is 0.45 and LCL 0.50.

TABLE II
MEASURED ON-DEVICE COMPARISON, STM32H7B3I-DK (CORTEX-M7, 280 MHZ), ST EDGE AI CORE 4.0.1, FP32.

	Reference CNN	Ours	Ours (tiny)
Parameters	441 k	84 k	8 k
Accuracy (%)	91.7	92.1	91.3
Latency (ms)	33.52	3.63	0.79
Flash (KB)	1729	335	37
RAM (KB)	38.3	9.2	9.2
MACC (M)	1.97	0.16	0.03

parameters (our RMSE 0.45 LCR, 0.47 LCL). We do not claim to beat the internal-reference Transformers on regression RMSE (0.42/0.44), which remain competitive small models; an RMSE-aware search narrowed but did not close that gap. Removing the turn-signal channels drops classification accuracy by about one point (92.1% to 91.1%)—evidence that the model anticipates from vehicle dynamics and surrounding traffic rather than reading a declared signal. [TODO: RMSE columns and the internal-reference RMSE (0.42/0.44) are not yet matched because the search optimizes MAE; decide RMSE-aware objective vs MAE-primary framing.]

C. Deployment on STM32

Target platform: STM32H7B3I-DK discovery kit (Cortex-M7 at 280 MHz, 1.4 MB SRAM, 2 MB flash) — the same MCU family used by the baseline framework [1], which deployed FP32 models only. We therefore compare in FP32, like for like, benchmarking both models on the same board through the ST Edge AI Developer Cloud (Core 4.0.1, balanced optimization). Table II reports measured, not estimated, figures.

The searched model is $9.2 \times$ faster and $5.2 \times$ smaller than the reference CNN while being more accurate; a second operating point on the same front gives up 0.4 points of accuracy for a $42 \times$ speed-up and a $47 \times$ smaller model that infers in under a millisecond.

The smaller operating point also changes what hardware is reachable, and here the difference is categorical rather than incremental. On a NUCLEO-F401RE (Cortex-M4, 84 MHz, 512 KB flash) the tiny classifier runs in 4.376 ms and occupies 7.2% of flash, while the reference CNN does not fit at all—its 1729 KB of weights exceed the board’s entire flash by

3.4 $\times$, which no compiler setting can recover. Nor is the small model the only one that reaches this board: the full classifier runs there in 18.35 ms at 65.5% of flash, and its quantized form in 7.381 ms at 25.6%, so the accuracy headline itself is available on the Cortex-M4, as is the LCL regression model at 162.5 ms. The searched models therefore open a device class the reference is locked out of, consistent with the baseline’s own report that its Transformer did not fit this board. One build does not follow: the widest regressor in float32 occupies 90.5% of flash and returned no measured time, while every build that ran leaves at least 100 KB free. Quantizing that same network to int8—identical architecture, operators and board, 28.7% of flash—makes it run in 28.10 ms, which isolates the cause: the binding constraint is flash *headroom* for the runtime, not model size or operator support. Size arithmetic alone would have called the float32 build deployable, and on this board quantization is therefore not only a speed and memory optimisation but the difference between a model that runs and one that does not.

One measurement is worth recording for anyone budgeting from operation counts: the tiny model costs 7.0 cycles per MAC on the M7 and 11.6 on the M4, on cores with single-cycle MAC instructions. At this scale inference is bound by per-op overhead and memory traffic, not arithmetic, and per-layer time does not follow per-layer MACs—a zero-MAC pooling operator is among the most expensive layers in both classifiers. MAC count remains a good predictor of whole-model latency across models, but a poor one within a model.

The per-layer breakdown explains why. Both our classifier and the reference spend most of their flash in the fully-connected head, so the gain is not from avoiding a large head but from what enters it: the reference flattens a 50×128 map into 6400 features, giving a 64×6400 layer of 1.64 MB (93% of its weights), whereas the searched models pool first and flatten only 308 (or 98) features, cutting that layer 5.9 $\times$.

Across the four searched models this reliance falls monotonically with how much of the search space the task exploits. The largest fully-connected layer holds 93% of weights in the reference and 90.9% in our tiny classifier, but 39.6% in the LCR model and only 11.4% in the LCL model—whose largest layer is a convolution. The same ordering appears in topology: both classifiers are essentially sequential, whereas the LCR and LCL models fan out into two and three parallel branches respectively. The search space permits branches for every task; only regression used them. Topology follows the task, discovered rather than imposed.

RAM behaves differently and is worth stating plainly, because three distinct regimes appear across our models. Both searched classifiers measure ~ 9.2 KB regardless of a $9 \times$ parameter gap: the 50×31 float32 input tensor alone occupies 6.2 KB and no architecture can go below it, so they are *input-bound*. The LCR model measures 20.8 KB and is *width-bound*: its widest branch emits a 25×116 tensor of 11.6 KB, and the peak working set is that plus the input. The LCL model measures 26.5 KB and is neither—it is *liveness-bound*. Its widest activation (13.0 KB) and the next-widest (12.0 KB) have disjoint live ranges and are never co-resident, so no width-based rule reaches the measurement; instead its three-

way merge pins the 6.2 KB input and two 3.3 KB branch outputs across many operators, giving a 24.8 KB peak, with the remainder attributable to allocator fragmentation. The general statement is therefore that peak RAM is the maximum over the schedule of all simultaneously live tensors, which collapses to $\max(\text{input}, \text{widest in} + \text{out})$ only for a chain. This is why the RAM advantage ($\sim 4 \times$) is far smaller than the flash advantage.

The input-bound regime is not a property of the network but of the tensor interface, and we verified this by changing only the interface. Exporting the int8 classifier with int8 rather than float32 input and output tensors—the weights, graph and accuracy are untouched, 86.86% either way—removes the 6.2 KB float32 input buffer and both cast operators, and measured RAM falls from 8,096 to 5,444 B. The reduction is smaller than the 4.65 KB of buffer removed, which is the same maximum-over-live-tensors point again: shrinking the input does not subtract from the arena, it merely stops the input being the binding constraint, after which the widest internal activation sets the floor at 3.5 KB. Reporting int8 with a float32 interface therefore understates its memory advantage—1.17 $\times$ against float32 before the change, 1.74 $\times$ after—and we note it because the default converter settings produce exactly that misleading configuration.

On quantization, the ST toolchain accepts only float32/int8/uint8 tensors. Full post-training int8 shrinks the classifier to 104 KB and 1.885 ms but costs 5.2 accuracy points on these wide-dynamic-range inputs (92.1% $\rightarrow$ 86.9%; LCR MAE 0.287 $\rightarrow$ 0.449), so we report float32 as the headline configuration—which also matches the baseline’s own FP32 deployment. Quantization-aware fine-tuning recovers most of that int8 loss: the classifier returns to 89.8% (from 86.9% post-training), closing 57% of the gap to float32, and on the board it runs in 1.558 ms—the fastest of the three operating points—in 128 KB, still 2.6 $\times$ below the float32 build. The tooling quantizes only 2D operators, so the 1D graph is deployed as width-1 2D convolutions; the computation is identical and the accuracy is read through the int8 interpreter. The same procedure does *not* transfer to the regression heads: it recovers only 11% of the gap for LCR and degrades LCL outright (0.344 $\rightarrow$ 0.362 s MAE), a result stable across fine-tuning rates. We read this as a task asymmetry rather than a tuning failure—classification needs only the argmax of three logits, which tolerates coarse activations, while a continuous time-to-lane-change estimate does not, and int8 costs the classifier 5.7% of its accuracy against a 57% increase in LCR error. Weight fine-tuning cannot restore information lost in the activation representation, so we report float32 as the only accurate regression configuration and int8 there purely as a deployability fallback. An int16 $\times$ 8 scheme (int8 weights, int16 activations) preserves accuracy almost exactly offline (92.1% $\rightarrow$ 92.2%; LCR MAE 0.287 $\rightarrow$ 0.286), but the toolchain does not deploy it. We verified this on hardware rather than trusting the documentation: a 117.9 KB int16 $\times$ 8 file returns 326.15 KB of weights, byte-identical to the float32 build, while a genuinely int8 control compresses to 83.28 KB. The scheme is silently dequantized to float32, so we report it as an offline bound only. The episode is worth one sentence of method: a toolchain that accepts an

unsupported scheme and quietly widens it will produce a benchmark that looks like a result. [TODO: cite ST Edge AI Core 4.0 documentation (quantization / supported ops, retrieved 2026-07-14).]

VI. CONCLUSION

[TODO: two or three sentences; no new claims.]

REFERENCES

- [1] L. Forneris *et al.*, “A deployment-oriented simulation framework for deep learning-based lane change prediction,” *IEEE Signal Process. Lett.*, vol. 33, pp. 136–140, 2026, doi: 10.1109/LSP.2025.3638676.
- [2] E. Liberis, Ł. Dudziak, and N. D. Lane, “ μ NAS: Constrained neural architecture search for microcontrollers,” in *Proc. 1st Workshop Mach. Learn. Syst. (EuroMLSys)*, 2021, pp. 70–79.
- [3] L. Forneris *et al.*, “Lane change intention recognition dataset,” Zenodo, 2025, doi: 10.5281/zenodo.16686054.
- [4] L. Forneris *et al.*, “Setting up a lightweight simulation environment for automated driving dataset collection,” in *Applications in Electronics Pervading Industry, Environment and Society (ApplePies)*, Springer LNEE, 2024, pp. 156–163, doi: 10.1007/978-3-031-84100-2_19.
- [5] B. Zoph and Q. V. Le, “Neural architecture search with reinforcement learning,” in *Proc. Int. Conf. Learn. Representations (ICLR)*, 2017.
- [6] E. Real, A. Aggarwal, Y. Huang, and Q. V. Le, “Regularized evolution for image classifier architecture search,” in *Proc. AAAI Conf. Artif. Intell.*, 2019, pp. 4780–4789.